

# NAME

`zprof`, `zcounts` - transform ANSI-C programs to enable basic block profiling.

# SYNOPSIS

```
zprof
  [--cond|-c]
  [--help|-h]
  [--prefix=PREFIX|-pPREFIX]
  [--preprocess|-P]
  [--silent|-s]
  [--version|-v]
  [--typedef=ID|-tID]
[file.c...]

zcounts
```

# DESCRIPTION

**zprof** is a preprocessor which transforms a ANSI-C program to enable basic block profiling. When the transformed program is compiled and executed, it records the number of times each of its basic blocks is executed.

Each source file *file.c* is transformed to *file\_BB.c* (to change the output name, see the **--prefix** option). The transformed program needs to be linked with the **zprof** library **libzprof.a**. When the transformed program is executed, it appends the execution counts of its basic blocks to the file **zprof.out**. The **zcounts** script uses the generated **zprof.out** file to produce an annotated version of the original files (with names of the form *file.c.bb*) with each line preceded by its execution count. The **zprof.out** file can also be analyzed using standard utilities (see the **EXAMPLES** section).

If no source files are specified, **zprof** simply prints a summary of its options and exits.

# OPTIONS

**Zprof** supports both traditional single-letter options and mnemonic long option names. Long option names are indicated using `--` instead of `-`. Abbreviations for long option names are allowed as long as they are unique. When a long option like `--prefix` takes an argument, connect the option name and the argument with `=` or put the argument in the next word. Similarly, when a short option like `-p` takes an argument, put the argument value immediately after the option in the same word, or put the argument in the next word. Boolean argument values may be specified as `0` or `1`, `yes` or `no`, `on` or `off`.

**--cond**

**-c** Profile subexpressions of all `?:` conditional expressions (default: `off`).

**--help**

**-h** Print summary of options and exit.

**--prefix=PREFIX**

**-pPREFIX**

Use `PREFIX` as prefix of all generated names in instrumented file and for forming the name of the output files (default: `_BB`).

**--preprocess**

**-P** Run the preprocessor (given by environmental var `CPP`) on the input file before **zprof** sees it (default: `on`).

**--silent**

**-s** Do not generate error messages (default: `off`).

**--version**

**-v** Print version information and exit.

**--typedef=ID**

**-tID** Declare identifier `ID` to be a typedef in the global scope.

# ENVIRONMENT

If preprocessing is not turned off using `--preprocess=no`, then the command used to run the C-preprocessor is taken from the environmental variable `CPP`. If `CPP` is not defined, then the name of a C compiler is sought in environmental variable `CC`, and each source file `file.c` is preprocessed using the command:

```
$CC -E file.c
```

If **CC** is also not defined, then the C-preprocessor is run using

```
cc -E file.c
```

## EXAMPLES

Consider the simple prime number program from file **primes.c** shown below:

```
static int
prime(int n)
{
    int i;
    for (i= 2; i < n; i++) {
        if (n % i == 0) {
            return 0;
        }
    }
    return 1;
}

int
main()
{
    const int n= 1000;
    int i;
    for (i= 2; i <= n; i++) {
        if (prime(i)) {
            printf("%d0, i);
        }
    }
    return 0;
}
```

This can be profiled as follows:

```
$ zprof primes.c
$ cc primes_BB.c -lzprof -o primes_bb
$ rm -f zprof.out
$ ./primes_bb >/dev/null
$ zcounts
```

The first line transforms **primes.c** into **primes\_BB.c** which has additional code for accumulating profile counts. The second line compiles the transformed file and links the object file with the **zprof** library to get an executable **primes\_bb**. The third line ensures that there is no **zprof.out** file present in the current directory, as the execution counts from the next step are appended to it. The fourth line runs **primes\_bb**, producing a **zprof.out** file

containing the execution statistics. The final line uses the statistics to produce an annotated version **primes.c.bb** shown below:

```

static int
prime(int n)
{
    int i;
999    for (i= 2; i < n; i++) {
78022        if (n % i == 0) {
831            return 0;
        }
    }
168    return 1;
}

int
main()
{
    const int n= 1000;
    int i;
1    for (i= 2; i <= n; i++) {
999        if (prime(i)) {
168            printf("%d0, i);
        }
    }
1    return 0;
}

```

In order to correctly parse C, it is necessary to recognize **typedefs**. Since **typedefs** are often defined in header files, it is necessary to parse all included header files --- this is typically done by running the C-preprocessor first. Unfortunately, on some systems like Linux, the system header files use non-ANSI extensions, which causes problems for **zprof** which expects ANSI-C. It is still possible to use **zprof** on such systems by using the C-preprocessor to remove the non-ANSI extensions. For example, the following works for profiling some programs under Linux.

```

$ CPP="gcc -E -D'__attribute(n)___=' -D__const=const" \
zprof ...

```

The options **--preprocess=0**, **--silent** and **--typedef** may also be useful in such situations where a header which cannot otherwise be parsed defines only a few **typedefs**.

## FILES

The format of the generated **zprof.out** file consists of lines of the form:

```
filename:linenum: count
```

specifying that line number *linenum* in file *filename* was executed *count* times. If there are multiple basic blocks associated with a particular source line, then multiple lines will be produced for the same *filename* and *linenum*.

The above format is amenable to easy processing by several tools:

- + The file can easily be sorted in descending order of counts:

```
sort -t: -nr -2 zprof.out
```

- + Emacs compilation-mode can be used to interactively browse this file. **M-x compile** can be used to start compilation and the compilation command can be specified simply as **cat zprof.out**. Then the **M-`** command can be used repeatedly to position the cursor at each source line with its execution count shown in the compilation buffer.

## AUTHOR

Copyright (C) 1997 Zerksis D. Umrigar

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

## SEE ALSO

**prof**(1) **monitor**(3)